

LLVM Intermediate Representation

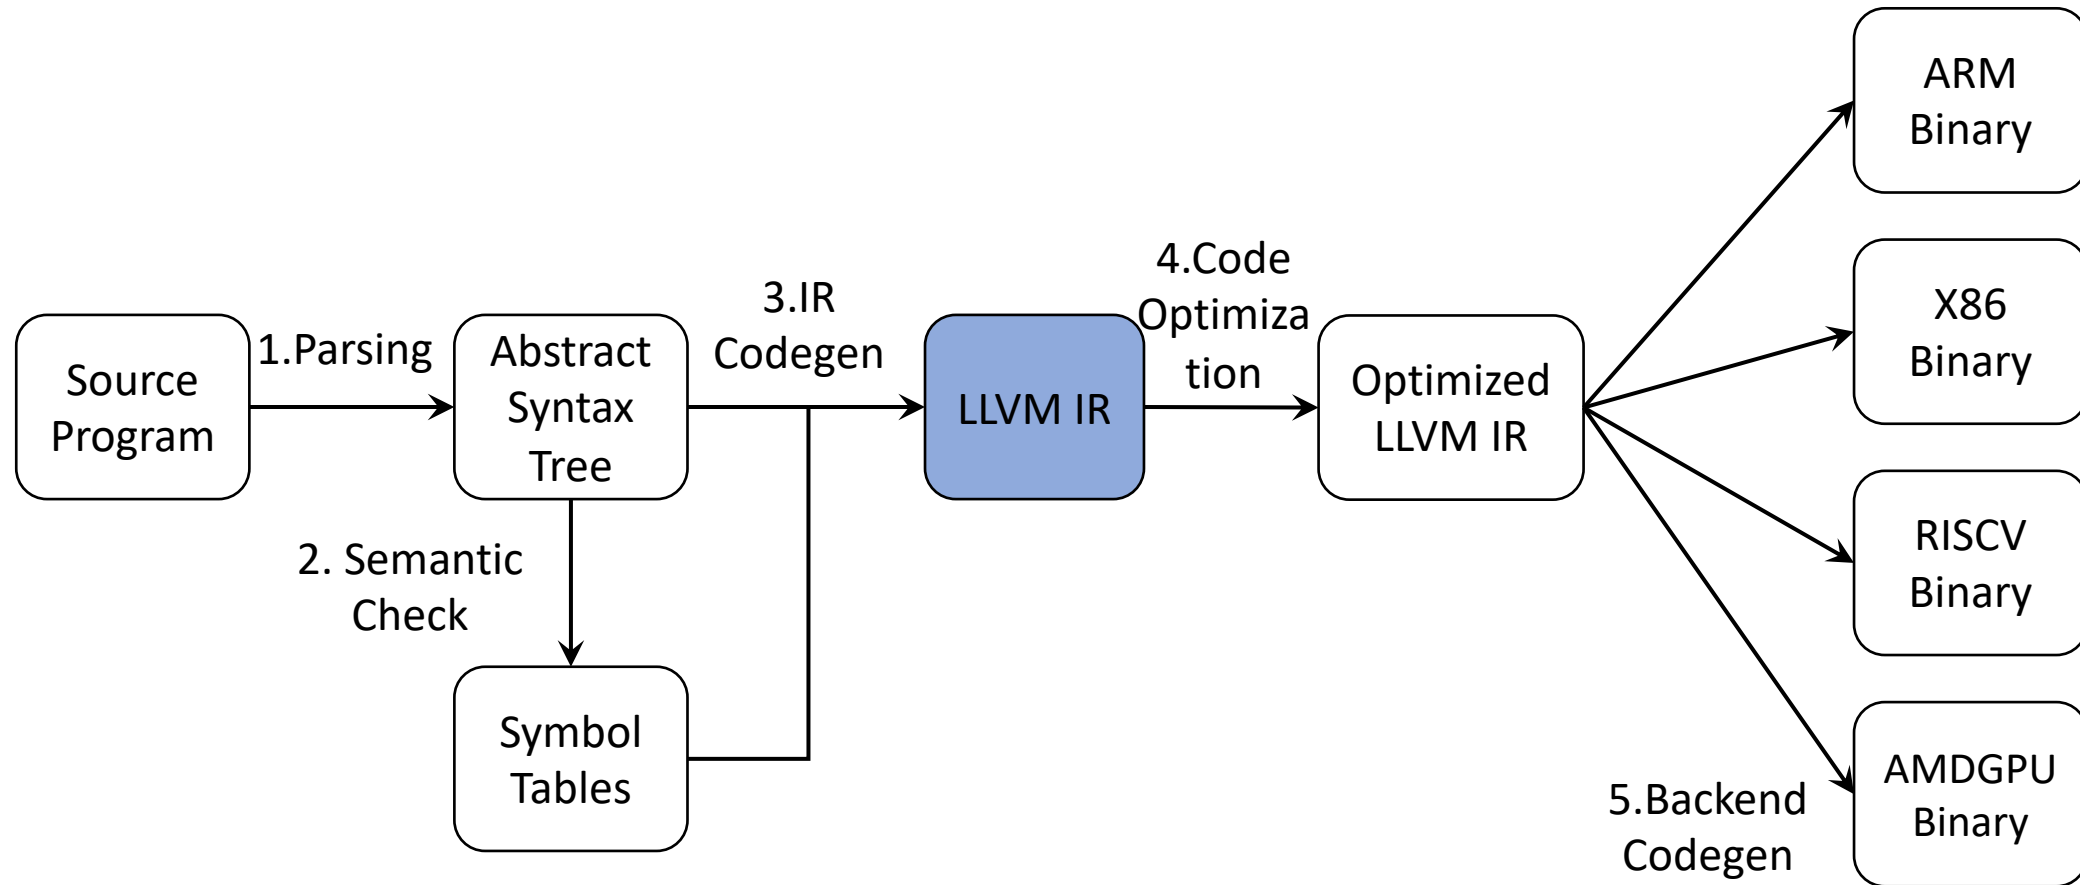
Fan Long

University of Toronto

What is LLVM IR?

- LLVM IR stands for low-level virtual machine intermediate representation.
- An **universal** and **architecture-independent** IR for compiler optimization, code generation, and program analysis.
- All LLVM family compilers first compile source programs into LLVM IR, then perform optimizations and target code generations.

Typical LLVM Compiler Workflow



...

Advantages of Using LLVM IR

- Utilize existing code optimization passes in LLVM framework to generate **fast** code.
- Utilize existing backends in LLVM framework to generate binary code for **different architectures**, e.g., x86, ARM, RISC-V, etc.
- Utilize existing program analysis tools built on LLVM IR for your compiler.



LLVM IR At a Glance

C program language

LLVM IR

- | <i>C program language</i> | <i>LLVM IR</i> |
|--|--|
| • Scope: <i>file, function</i> | <i>module, function</i> |
| • Type: <i>bool, char, int, struct{int, char}</i> | <i>i1, i8, i32, {i32, i8}</i> |
| • A statement with multiple expressions | A sequence of instructions each of which is in a form of “x = y op z”. |
| • Data-flow:
a sequence of reads/writes on variables | <ol style="list-style-type: none">1. load the values of memory addresses (variables) to registers;2. compute the values in registers;3. store the values of registers to memory addresses <p>* each register must be assigned exactly once (SSA)</p> |
| • Control-flow in a function:
if, for, while, do while, switch-case,... | A set of basic blocks each of which ends with a conditional jump (or return) |

LLVM IR Example

simple.c

```
1  #include <stdio.h>
2  int x, y ;
3
4  int main() {
5      int t ;
6      scanf("%d %d",&x,&y);
7      t = x - y ;
8      if (t > 0)
9          printf("x > y") ;
10     return 0 ;
11 }
```

```
$ clang -S -emit-llvm simple.c
```

simple.ll (simplified)

```
...
2  6 @x = common global i32 0, align 4
   7 @y = common global i32 0, align 4
...
4  11 define i32 @main() #0 {
   12     entry:
...
5  14 %t = alloca i32, align 4
...
6  16 %call = call i32 @__isoc99_scanf(...i32* @x,i32* @y)
...
7  17 %0 = load i32* @x, align 4
   18 %1 = load i32* @y, align 4
   19 %sub = sub nsw i32 %0, %1
   20 store i32 %sub, i32* %t, align 4
...
8  21 %2 = load i32* %t, align 4
   22 %cmp = icmp sgt i32 %2, 0
   23 br i1 %cmp, label %if.then,
        label %if.end
...
9  24 if.then:
   25     %call1 = call i32 @printf(...)
   26     br label %if.end
...
10 27 if.end:
   28     ret i32 0
```

Content

- LLVM IR Instruction
 - architecture, static single assignment
- Data Representation
 - types, constants, registers, variables
 - load/store instructions, cast instructions
 - computational instructions
- Control Representation
 - control flow (basic block)
 - control instructions
- How to generate LLVM IR?

* *LLVM Language Reference Manual* <http://llvm.org/docs/LangRef.html>

* *Mapping High-Level Constructs to LLVM IR*

<https://mapping-high-level-constructs-to-llvm-ir.readthedocs.io/en/latest/>

LLVM IR Architecture

- RISC-like instruction set
 - Only 31 op-codes (types of instructions) exist
 - Most instructions (e.g., computational instructions) are in three-address form: one or two operands, and one result
- Load/store architecture
 - Memory can be accessed via load/store instruction
 - Computational instructions operate on registers
- Infinite and typed *virtual registers*
 - It is possible to declare a new register any point (the backend maps virtual registers to physical ones).
 - A register is declared with a primitive type (boolean, int, float, pointer)

Static Single Assignment (SSA)

- In SSA, each variable is assigned exactly once, and every variable is defined before its uses.
- Conversion
 - For each definition, create a new version of the target variable (left-hand side) and replace the target variable with the new variable.
 - For each use, replace the original referred variable with the versioned variable reaching the use point.

```
1  x = y + x ;
2  y = x + y ;
3  if (y > 0)
4      x = y ;
5  else
6      x = y + 1 ;
```



```
11 x1 = y0 + x0 ;
12 y1 = x1 + y0 ;
13 if (y1 > 0)
14     x2 = y1 ;
15 else
16     x3 = y1 + 1 ;
```


Static Single Assignment (SSA)

- LLVM IR follows the SSA form.
 - Every virtual register can be only assigned once in the program.
 - Memory loads/stores are not affected by this rule.
 - Multiple assignments to a same variable have to be mapped to multiple virtual registers.
 - Therefore, `llvm::Value` is the base class of all LLVM IR elements.
 - Every instruction is uniquely associated with the value it defines (assigns to).
- How to handle local variables being modified in different branches?
 - **Option 1:** Use **`alloca`** instruction to allocate stack memory space to hold local variables and then use memory loads/stores
 - **Option 2:** Use **`phi`** instructions

SSA and Phi Functions

- Use ϕ function if two versions of a variable are reaching one use point at a joining basic block
 - $\phi(x_1, x_2)$ returns either x_1 or x_2 depending on which block was executed

```
1  x = y + x ;
2  y = x + y ;
3  if (y > 0)
4      x = y ;
5  else
6      x = y + 1 ;
7  y = x - y ;
```



```
11 x1 = y0 + x0 ;
12 y1 = x1 + y0 ;
13 if (y1 > 0)
14     x2 = y1 ;
15 else
16     x3 = y1 + 1 ;
17 x4 =  $\phi(x2, x3)$  ;
18 y2 = x4 - y1 ;
```

Data Representation

- Primitive types
- Constants
- Registers (virtual registers)
- Variables
 - local variables, heap variables, global variables
- Load and store instructions
- Aggregated types

Primitive Types

- Language independent primitive types with predefined sizes
 - void: **void**
 - bool: **i1**
 - integers: **i[N]** where **N** is 1 to $2^{23}-1$
e.g. **i8**, **i16**, **i32**, **i1942652**
 - floating-point types:
 - half** (16-bit floating point value)
 - float** (32-bit floating point value)
 - double** (64-bit floating point value)
- Pointer type is a form of **<type>*** (e.g. **i32***, **(i32*)***)

Constants

- Boolean (i1): **true** and **false**
- Integer: standard integers including negative numbers
- Floating point: decimal notation, exponential notation, or hexadecimal notation (IEEE754 Std.)
- Pointer: **null** is treated as a special value

Registers

- Identifier syntax
 - Named registers: `[%] [a-zA-Z$. _] [a-zA-Z$. _0-9] *`
 - Unnamed registers: `[%] [0-9] [0-9] *`
- A register has a function-level scope.
 - Two registers in different functions may have the same identifier
- A register is assigned for a particular type and a value at its first (and the only) definition (**SSA form**)

Variables

- In LLVM, all addressable objects (“lvalues”) are explicitly allocated.
- Global variables
 - Each variable has a global scope symbol that points to the memory address of the object
 - Variable identifier: `[@] [a-zA-Z$. _] [a-zA-Z$. _0-9] *`
- Local variables
 - The **alloca** instruction allocates memory in the stack frame.
 - Deallocated automatically if the function returns.
- Heap variables
 - The **malloc** function call allocates memory on the heap.
 - The **free** function call frees the memory allocated by **malloc**.

Load and Store Instructions

- Load


`<result>=load <type>* <ptr>`

- result: the target register
- type: the type of the data (a pointer type)
- ptr: the register that has the address of the data

- Store

`store <type> <value>,<type>* <ptr>`

- type: the type of the value
- value: either a constant or a register that holds the value
- ptr: the register that has the address where the data should be stored



Variable Example

```
1 #include <stdlib.h>
2
3 int g = 0 ;
4
5 int main() {
6     int t = 0;
7     int * p;
8     p=malloc(sizeof(int));
9     free(p);
10 }
```



```
1 @g = global i32 0, align 4
...
8 define i32 @main() #0 {
...
10 %t = alloca i32, align 4
11 store i32 0, i32* %t, align 4
12 %p = alloca i32*, align 8
13 %call = call noalias i8*
    @malloc(i64 4) #2
14 %0 = bitcast i8* %call to i32*
15 store i32* %0, i32** %p,
    align 8
16 %1 = load i32** %p, align 8
...

```

Aggregate Types and Function Type

- Array: **[<# of elements> x <type>]**
 - Single dimensional array ex: `[40 x i32], [4 x i8]`
 - Multi dimensional array ex: `[3 x [4 x i8]], [12 x [10 x float]]`
- Structure: **type {<a list of types>}**
 - E.g., `type{ i32, i32, i32 }, type{ i8, i32 }`
- Function: **<return type> (a list of parameter types)**
 - E.g., `i32 (i32), float (i16, i32*) *`

Getelementptr Instruction

- A memory in an aggregate type variable can be accessed by **load/store** instruction and **getelementptr** instruction that obtains the pointer to the element.

- Syntax:

<res> = getelementptr <pty>* <ptrval>{,<t> <idx>}*

- res: the target register
- pty: the register that defines the aggregate type
- ptrval: the register that points to the data variable
- t: the type of index
- idx: the index value

Aggregate Type Example 1

```
1 struct pair {  
2     int first;  
3     int second;  
4 };
```

```
5 int main() {  
6     int arr[10];  
7     struct pair a;
```

```
8     a.first = arr[1];
```

...

```
11 %struct.pair = type{ i32, i32 }
```

```
12 define i32 @main() {  
13     entry:
```

```
14     %arr = alloca [10 x i32]
```

```
15     %a = alloca %struct.pair
```

```
16     %arrayidx = getelementptr  
        [10 x 32]* %arr, i32 0, i64 1
```

```
17     %0 = load i32*, %arrayidx
```

```
18     %first = getelementptr  
        %struct.pair* %a, i32 0, i32 0
```

```
19     %store i32 %0, i32* %first
```

Aggregate Type Example 2

```
1 struct RT {  
2     char A;  
3     int B[10][20];  
4     char C;  
5 };  
6 struct ST {  
7     int X;  
8     double Y;  
9     struct RT Z;  
10 };  
11  
12 int *foo(struct ST *s) {  
13     return &s[1].Z.B[5][13];  
14 }
```

```
5 %struct.RT = type { i8, [10 x [20 x i32]  
    ], i8 }  
6 %struct.ST = type { i32, double, %struct  
    .RT }  
7  
8 define i32* @foo(%struct.ST* %s)  
    nounwind uwtable readnone optsize  
    ssp {  
9 entry:  
10     %arrayidx = getelementptr inbounds  
        %struct.ST* %s, i64 1, i32 2,  
        i32 1, i64 5,  
        i64 13  
11     ret i32* %arrayidx  
12 }
```

Integer Conversion

- Truncate

- Syntax: `<res> = trunc <iN1> <value> to <iN2>` where `iN1` and `iN2` are of integer type, and `N1 > N2`

- Examples

- `%X = trunc i32 257 to i8 ; %X becomes i8:1`
- `%Y = trunc i32 123 to i1 ; %Y becomes i1:true`
- `%Z = trunc i32 122 to i1 ; %Z becomes i1:false`

Integer Conversion

- Zero extension

- `<res> = zext <iN1> <value> to <iN2>` where `iN1` and `iN2` are of integer type, and `N1 < N2`
- Fill the remaining bits with zero
- Examples
 - `%X = zext i32 257 to i64 ; %X becomes i64:257`
 - `%Y = zext i1 true to i32 ; %Y becomes i32:1`

- Sign extension

- `<res> = sext <iN1> <value> to <iN2>` where `iN1` and `iN2` are of integer type, and `N1 < N2`
- Fill the remaining bits with the sign bit (the highest order bit) of `value`
- Examples
 - `%X = sext i8 -1 to i16 ; %X becomes i16:65535`
 - `%Y = sext i1 true to i32 ; %Y becomes i32:232-1`

Other Conversions

- Float-to-float
 - `fptrunc .. to, fpext .. to`
- Float-to-integer (vice versa)
 - `fptoui .. to, tptosi .. to, uitofp .. to, sitofp .. to`
- Pointer-to-integer
 - `ptrtoint .. to, inttoptr .. to`
- Bitcast
 - `<res> = bitcast <t1> <value> to <t2>`
where `t1` and `t2` should be different types and have the same size

Computational Instructions

- Binary operations:
 - Add: `add`, `sub`, `fsub`
 - Multiplication: `mul`, `fmul`
 - Division: `udiv`, `sdiv`, `fdiv`
 - Remainder: `urem`, `srem`, `frem`
- Bitwise binary operations
 - shift operations: `shl`, `lshl`, `ashr`
 - logical operations: `and`, `or`, `xor`

Add Instruction Example

- `<res> = add [nuw][nsw] <iN> <op1>, <op2>`
 - nuw (no unsigned wrap): if unsigned overflow occurs, the result value becomes a poison value (undefined)
 - E.g: `add nuw i8 255, i8 1`
 - nsw (no signed wrap): if signed overflow occurs, the result value becomes a poison value
 - E.g. `add nsw i8 127, i8 1`

Control Flow Representation

- The LLVM front-end constructs the control flow graph (CFG) of every function explicitly in LLVM IR
 - A function has a set of **basic blocks** each of which is a sequence of instructions
 - A function has exactly **one entry basic block**
 - Every basic block is ended with exactly one **terminator instruction** which explicitly specifies its successor basic blocks if there exist.
 - Terminator instructions: branches (conditional, unconditional), return, unwind, invoke
- Due to its simple control flow structure, it is convenient to analyze, transform the target program in LLVM IR


Label, Return, and Unconditional Branch

- A label is located at the start of a basic block
 - Each basic block is addressed as the start label
 - A label `x` is referenced as register `%x` whose type is label
 - The label of the entry block of a function is “`entry`”
- Return `ret <type> <value> | ret void`
- Unconditional branch `br label <dest>`
 - At the end of a basic block, this instruction makes a transition to the basic block starting with label `<dest>`
 - E.g: `br label %entry`

Conditional Branch

- `<res> = icmp <cmp> <ty> <op1>, <op2>`
 - Returns either true or false (i1) based on comparison of two variables (op1 and op2) of the same type (ty)
 - `cmp`: comparison option
 - eq (equal), ne (not equal), ugt (unsigned greater than),
 - uge (unsigned greater or equal), ult (unsigned less than),
 - ule (unsigned less or equal), sgt (signed greater than),
 - sge (signed greater or equal), slt (signed less than), sle (signed less or equal)
- `br i1 <cond>, label <thenbb>, label <elsebb>`
 - Causes the current execution to transfer to the basic block <thenbb> if the value of <cond> is true; to the basic block <elsebb> otherwise.

```
1  if (x > y)
2      return 1 ;
3  return 0 ;
```




```
11 %0 = load i32*, %x
12 %1 = load i32*, %y
13 %cmp = icmp sgt i32 %0, %1
14 br i1 %cmp, label %if.then, label %if.end
15 if.then:
    ...
```

Switch

- **switch** <iN> <value>, label <defaultdest>
[<iN> <val>, label <dest> ...]
- Transfer control flow to one of many possible destinations
- If the value is found (val), control flow is transferred to the corresponding destination (dest); or to the default destination (defaultdest)

```
1  switch(x) {  
2      case 1:  
3          break ;  
4      case 2:  
5          break ;  
6      default:  
7          break ;  
8  }
```



```
11  %0 = load i32*, %x  
12  switch i32 %0, label %sw.default [  
13      i32 1, label %sw.bb  
14      i32 2, label %sw.bb1]  
15  sw.bb:  
16      br label %sw.epilog  
17  sw.bb1:  
18      br label %sw.epilog  
19  sw.default:  
20      br label %sw.epilog  
21  sw.epilog:  
    ...
```

PHI (Φ) instruction

- `<res> = phi <t> [<val_0>, <label_0>],`
`[<val_1>, <label_1>], ...`
 - Return a value `val_i` of type `t` such that the basic block executed right before the current one is of `label_i`

```

1  y = (x > 0) ? x : 0 ;

```

➡

```

11 %0 = load i32*, %x
12 %c = icmp sgt i32 %0, 0
13 br i1 %c, label %c.t, %c.f

14 c.t:
15 %1 = load i32*, %x
16 br label %c.end

17 c.f:
18 br label %c.end

19 c.end:
20 %cond = phi i32 [%1, %c.t], [0, %c.f]
21 store i32 %cond, i32* %y

```


Function Call

- `<res> = call <t> [<fnty>*] <fnptrval>(<fn args>)`
 - `t`: the type of the call return value
 - `fnty`: the signature of the pointer to the target function (optional)
 - `fnptrval`: an LLVM value containing a pointer to a target function
 - `fn args`: argument list whose types match the function signature

```
1  printf("%d", abs(x));
```



```
11 @.str = [3 x i8] c"%d\00"
12 %0 = load i32* %x
13 %call = call i32 @abs(i32 %0)

14 %call1 = call i32 (i8*, ...) *
    @printf(i8*
    getelementptr ([3 x i8]* @.str,
    i32 0, i32 0),
    i32 %call)
```

How to Generate LLVM IR?

- Option 1: Directly generate LLVM IR as texts following the syntax
 - Quick to get started
 - Lack syntax checking and verification
 - Lack semantic checking
- Option 2: Use LLVM framework API to build llvm modules
 - Built-in syntax and semantic checking
 - Recommended way for building serious compilers

How to Generate LLVM IR?

- Create a new Module object (`llvm::Module`).
- Create a Builder object associated with the module (`llvm::IRBuilder`).
- Create global variable (`llvm::GlobalVariable`) objects for the module.
- Create new Function objects (`llvm::Function`) for the Module.
- Create the entry basic block for the function (`llvm::BasicBlock`).
- Set the insertion point of the builder to the entry basic block.
- Call `CreateXXX()` methods in `IRBuilder` to insert new instructions.
- Create additional basic blocks and change the insertion point of the builder when needed.

Q/A?